

The Dark Side of Unikernels for Machine Learning

Matthew Leon
Vanderbilt University
Nashville, Tennessee, USA
matthew.leon@vanderbilt.edu

Abstract—This paper analyzes the shortcomings of unikernels as a method of deployment for machine learning inferencing applications as well as provides insights and analysis on future work in this space. The findings of this paper advocate for a tool to enable management of dependent libraries in a unikernel to enable a more ergonomic build process as well as take advantage of the inherent security and performance benefits of unikernels.

Index Terms—unikernel, virtualization, xen, kernel samepage merging, docker, containerization, lightweight operating system, library operating system, cloud computing

I. INTRODUCTION

Virtualization technology is used in datacenters spanning the whole world to provide availability, scalability, and security to millions of client workloads. While virtualizing an entire computer and running everything, including the OS image, libraries and application code, on top is still popular, many alternative methods have emerged over the last decade, each with promises to increase security and improve resource utilization. There are numerous benefits to improving resource utilization of the host system. For users, fewer resources used by the host operating system means more resources available to application code. For the large corporations hosting public clouds, maximizing utilization on a small number of machines means lower costs, as datacenters consume massive amounts of power [1]. Recently, containerization technology has been adopted as a way to reduce the number of virtualization layers in the modern datacenter, with services like Docker providing benefits including easier deployment, ensuring consistency between the development and production environments, providing limits on resource usage, and sandboxing applications for better security. Containerization cuts down on costs by reducing duplication of the operating system—rather than running several different stacks all virtualized on top of a hypervisor, a user can run a single operating system and divide up its resources among several containers. Unikernel is yet another lightweight virtualization technology increasingly being adopted in cloud data centers.

II. WHAT ARE UNIKERNELS?

Unikernels, on the other hand, focus on the other side of the playing field from containers. With unikernels, the operating system is totally eliminated—the application code itself is augmented with the minimal set of code necessary to interface with the hypervisor and is then directly run as a bootable image on top of a hypervisor. The compactness of this system can result in numerous benefits over containers and

fully virtualized linux servers. In one study, boot times as low as 50ms were achieved, as well as lower memory usage and reduced latency due to zero-copy network implementation [2]. The significant benefits of the unikernels are discussed in the next section.

III. STUDY GOALS

This study aimed to analyze the state of a few different unikernels and their environments, comparing them to traditional methods of virtualization in terms of developer experience, performance, flexibility, security, and feasibility for adoption. Specifically, the study was conducted through a use case where we wanted to understand whether it was feasible to deploy a machine learning-trained image classification inference inside a Unikernel. To that end, we implemented an image classification API capable of receiving an image via HTTP and responding with an inference as to the contents of the image.

IV. REPORT ORGANIZATION

This report first outlines preliminary knowledge about the differences of unikernels, including major vendors of unikernel technology as well as an overview of the pertinent differences from ordinary virtualization solutions. The next section provides an overview of the work done in the process of evaluating the maturity of unikernels as a modern, lightweight alternative to containerization technology. Finally, the paper is concluded with an analysis of the hurdles that must be addressed before unikernels are sufficient for a modern deployment,

V. UNIKERNELS IN-DEPTH

A. What are Unikernels?

Unikernels, on the other hand, focus on the other side of the playing field from containers. With unikernels, the operating system is totally eliminated—the application code itself is augmented with the minimal set of code necessary to interface with the hypervisor and is then directly run as a bootable image on top of a hypervisor. The compactness of this system can result in numerous benefits over containers and fully virtualized linux servers. In one study, boot times as low as 50ms were achieved, as well as lower memory usage and reduced latency due to zero-copy network implementation [2]. The significant benefits of the unikernels are discussed in the next section.

Unikernels can be grouped into two distinct categories. Firstly are unikernels that function as a library operating system. OSs in this group, such as IncludeOS [3], HaLVM [4], and MirageOS [5], cannot run full executable programs, instead, they are written in and run code in an augmented runtime environment that implements operating system functions, such as I/O. The other group of unikernels, such as RumpRun [6], and Nanos [7], provide application binaries an entire POSIX-compatible runtime environment which can run arbitrary ELF executables. In addition to these runtime environments, several build, orchestration, and packaging tools are available, such as ops [7], Unikraft [8], and UniK [9]. This study investigates the feasibility and shortcomings of using these tools to deploy a deep neural network inference solution available via a web API.

B. Benefits of Unikernels

The single address space architecture of unikernels provides numerous benefits that are not achievable with conventional preemptive multitasking operating systems. Firstly, the total attack surface is much lower with a unikernel. Bratterud, Happe, and Duncan highlight a 92% reduction in total bytes of code in a running unikernel, which they translate to a 92% smaller attack surface [10]. The lack of a shell prevents an entire class of vulnerabilities, while a single address space allows for compile-time address space layout randomization, which is more performant than the runtime alternative. In addition to the security implications of a single address space, the removal of kernel space eliminates time spent in kernel space context switches as well as scheduling interrupts by the guest OS. Instead, scheduling and load balancing is handled entirely by the hypervisor.

In terms of load balancing itself, unikernels offer distinct benefits for web-related tasks, especially due to their startup time. The unikernel itself being the executable and thus not requiring file systems to be initialized as well as the small size the kernel code occupies means that the only boot step necessary is initializing the network interface. In a hypervisor environment, this allows the unikernel to be booted in response to an incoming request in time to handle that request. Such a fast boot time allows horizontal scaling with the granularity of individual requests. This instant availability enables applications such as fog deployment for IoT, which was investigated by Cozzolino, Ding, and Ott [11]. This work is further being applied at the same time as this research as infrastructure in smart city monitoring of ongoing road hazards [12].

VI. INSIGHTS FROM OUR STUDY

Supplementary source code materials and motivating examples for the following findings may be found at [13]. Many simple implementations of image classifications are available on GitHub, such as [14]. In the goal of evaluating the effectiveness of unikernels in different environments and implementations, three different machine learning frameworks were tested: Tensorflow, PyTorch, and Tensorflow.js (Tensorflow and Tensorflow.js are included separately as they do not

share bindings to the same underlying library; they are completely separate implementations in two different languages of the same API). IncludeOS was used in conjunction with Tensorflow, and RumpRun and Nanos were both used to test each of PyTorch and Tensorflow.js.

Our findings revealed that none of the tested solutions were successful. The shortcomings ranged depending on the implementation—Tensorflow and PyTorch struggled with issues linking inside of the unikernel, and Tensorflow.js struggled fetching the trained model via URL due to the lack of a DNS resolver in the unikernel environment. When adding the node.js extension to Tensorflow.js to allow for loading the model from within the image, the unikernel struggled due to lack of node-gyp (a C/C++ native binding) support inside the unikernel. We note that Tensorflow.js could be extended to support loading from file without involving node-gyp, but performing large modifications to the source of the application was out of scope for this study’s investigation of unikernels as an alternative deployment environment. PyTorch encountered similar issues as it is an optimized runtime with most of the deep learning code implemented in C—the modules for the library were unable to be loaded inside the unikernel environment.

Seeing as most of the encountered issues were due to the lack of interoperability between native libraries and interpreted code, the next approach we took was compiling Tensorflow into an application compiled with IncludeOS, the library operating system capable of transforming the C/C++ application it is built with to an Xen-bootable executable. Unfortunately, linking also became an issue in this case. The publicly available distributions of Tensorflow depend on over 10 shared libraries, and IncludeOS must be built statically, which is not supported (nor possible in an unsupported fashion) in any version of the library. Copying the shared libraries into the image from the system used to build Tensorflow resulted in a bootable system, but the execution failed due to missing symbols in the outdated version of glibc used in the host system. No other languages were tested after these failures, as all languages link to the C library, with the only exception being the previously mentioned Tensorflow.js without node.js extensions, which is designed for the browser environment. It was unexplored whether other smaller toolkits would’ve been more successful—mlpack [15] appears to be a good candidate for future research, as it may allow static linking [16].

VII. ANALYSIS AND POSSIBILITIES FOR FUTURE WORK

Unikernels, when compared to a deployment solution using docker containers or a native Linux virtual machine, still have many hurdles to overcome before they can claim full parity in terms of supported use cases. Due to the decades of prevalence of ecosystems which support dynamic linking as a way to quickly fix security issues and reduce compiled code duplication across binaries, even common libraries like Tensorflow do not support static linking, which is unfortunate news for any application developer looking to use these libraries in a unikernel. There are ways to build a static library manually

such as by packing GCC's object file output with tools such as ar, but these are steps for build system maintainers rather than application developers [17]. It is this researcher's opinion that unikernels would be most benefited by a robust build tool which handles dependency bundling inside the unikernel environment, much like Docker's *build* command or Ansible scripts. With access to a layered build system, unikernels could provide a compelling base layer for virtualization due to their lightweight and secure runtimes; however, dependencies in docker are handled through Linux distribution archives, which would be lacking in the environment of a unikernel. Without such tools, the art of manually packing a static archive for linking or building each shared library with the correct version of Glibc will remain out of reach for all but the most skilled devops engineers deploying in the most demanding situations where significant cost and performance benefits of unikernels may offset the additional development work required for deploying the unikernel. The build tools tested during this study, unik and ops, were both unable to contend with library dependencies in an efficient manner.

Beyond the deployment itself, there are supplementary considerations that must be investigated in terms of the performance implications of unikernels. Docker's AUFS allows for something which unikernel images, in their current, statically-linked form, do not—deduplication of layers. For example, if the unikernel is being used for a microservice-based web API, it would not be uncommon for there to be two endpoints that look very similar from a dependency point of view—endpoints involved with creating and updating a user's profile, for one—which would duplicate all library code in each binary. In a Docker deployment, the libraries for the operating system would be shared on disk, as the containers are stored as layers and extended with each command executed in the Dockerfile. This benefit extends to memory, as well—different docker containers descending from the same parent layers are able to share the same pages in memory due to Kernel Samepage Merging [18], [19]. Unikernels, on the other hand, may be able to share less memory due to differences in how private pages may be accessed by unikernels sharing a majority of code, but being compiled with different static dependencies. This is an area requiring further research to experimentally determine the extent of the memory saving, and the concept of copy-on-write deduplication of memory pages is currently subject to security concerns discovered along with side channel attacks [19] [20].

VIII. CONCLUSION

Unikernels present compelling benefits in terms of performance and security for deploying applications to the fog or the cloud, but currently face issues in regards to managing dependencies, updates, and compatibility with 3rd party libraries. A solution à la *docker build* for unikernels—providing a method for dependency management as well as possibly for sharing

and extending images others have made—may provide a more secure and performant platform for future cloud computing needs.

REFERENCES

- [1] D. E. Business. Powering a google search: The facts and figures. [Online]. Available: <https://business.directenergy.com/blog/2017/november/powering-a-google-search>
- [2] A. Madhavapeddy, R. Mortier, C. Rotsos, D. Scott, B. Singh, T. Gazagnaire, S. Smith, S. Hand, and J. Crowcroft, "Unikernels: Library operating systems for the cloud," *SIGARCH Comput. Archit. News*, vol. 41, no. 1, p. 461–472, Mar. 2013. [Online]. Available: <https://doi.org/10.1145/2490301.2451167>
- [3] N. Hussein, "IncludeOS: a unikernel for C++ applications," *LWN.net*. [Online]. Available: <https://lwn.net/Articles/728682/>
- [4] GaloisInc. The haskell lightweight virtual machine (HALVM) source archive. [Online]. Available: <https://github.com/GaloisInc/HaLVM>
- [5] Xen Foundation and Linux Foundation. (2017) MirageOS. [Online]. Available: <https://mirage.io/>
- [6] A. Kantee, "Flexible operating system internals: The design and implementation of the anykernel and rump kernels; flexible operating systems: Design and implementation of the kernel and stub kernels," p. 358, 2012. [Online]. Available: <http://urn.fi/URN:ISBN:978-952-60-4917-5>
- [7] NanoVMs Inc., "NanoVMs in depth," Tech. Rep., 2017. [Online]. Available: <https://nanovms.com/whitepapers>
- [8] NEC Laboratories Europe GmbH. (2020) Home — unikraft. [Online]. Available: <http://unikraft.org/>
- [9] Solo IO. The unikernel & microVM compilation and deployment platform. [Online]. Available: <https://github.com/solo-io/unik>
- [10] A. Bratterud, A. Happe, and R. Duncan, "Enhancing cloud security and privacy: The unikernel solution," in *Eighth International Conference on Cloud Computing, GRIDs, and Virtualization, 19 February 2017 - 23 February 2017, Athens, Greece*, ser. Cloud Computing IARIA. Curran Associates, 2 2017, pp. 79–86.
- [11] V. Cozzolino, A. Y. Ding, and J. Ott, "FADES: Fine-grained edge offloading with unikernels," in *Proceedings of the Workshop on Hot Topics in Container Networking and Networked Systems*, ser. HotConNet '17. New York, NY, USA: Association for Computing Machinery, 2017, p. 36–41. [Online]. Available: <https://doi.org/10.1145/3094405.3094412>
- [12] V. Cozzolino, J. Ott, A. Ding, and R. Mortier, "ECCO: Edge-cloud chaining and orchestration framework for road context assessment," in *Proceedings of the 2020 IEEE/ACM Fifth International Conference on Internet-of-Things Design and Implementation*, 2020.
- [13] M. Leon. Supplementary materials for "The Dark Side of Unikernels for Machine Learning". [Online]. Available: <https://github.com/leonm1/unikernel-research-2020>
- [14] avinassh. Pytorch flask API. [Online]. Available: <https://github.com/avinassh/pytorch-flask-api>
- [15] mlpack Contributors. (2020) mlpack. [Online]. Available: <https://github.com/mlpack/mlpack>
- [16] —. mlpack CMakeLists.txt. [Online]. Available: <https://github.com/mlpack/mlpack/blob/1c25a1bda52832841efec4e41477bfcbbbf6f3f/CMakeLists.txt#L30>
- [17] D. McKay. How to use linux's ar command to create static libraries. [Online]. Available: <https://www.howtogeek.com/427086/how-to-use-linuxs-ar-command-to-create-static-libraries/>
- [18] R. Waldon, M. Crosby, A. Suda, S. van Stijn, P. Gervai, and O. Veits. Question: does docker / LXC de-duplicate memory? if not, is it feasible? [Online]. Available: <https://github.com/moby/moby/issues/7950>
- [19] F. Shaikh, F. Yao, I. Gupta, and R. H. Campbell, "VMDedup: Memory de-duplication in hypervisor," in *2014 IEEE International Conference on Cloud Engineering*, 2014, pp. 379–384.
- [20] K. Suzuki, K. Iijima, T. Yagi, and C. Artho, "Software side channel attack on memory deduplication," *Proceedings of the 23rd ACM Symposium on Operating Systems Principles, Poster*, 01 2011.